

NLU course projects

Nicola Maestri (239920)

University of Trento

nicola.maestri@studenti.unitn.it

1. Introduction

In this assignment, we develop a *language model* (LM) for joint *intent classification* and *slot filling*. Initially, we implement an *LSTM*-based architecture, followed by separate classifiers for intent and slot predictions. Performance improves using a *bidirectional LSTM* and adding *dropout layers*. Integrating *BERT* as a foundation model and leveraging its contextual embeddings further improves test scores by over $> 2.5\%$ in both tasks. Finally, we evaluate the impact of fine-tuning all *BERT* layers, half of them and only the last one, observing a performance drop in *slot filling* for reduced fine-tuning.

2. Implementation details

In the first experiment, we use an *LSTM* followed by separate classifiers for intent and slot predictions. Additional enhancements include bidirectional *LSTMs* and dropout layers applied to embeddings and hidden states. Network hyperparameters are detailed in Table 1.

In the second experiment, *BERT* is adapted for intent and slot classification, as outlined in [1]. The *CLS* embedding is passed through a dropout layer and a linear classifier for intent classification. Similarly, the last hidden states are passed through a dropout layer and a linear classifier for slot filling. To address sub-tokenization, only the first sub-token of each word is used for slot classification. Slot sequences are padded and aligned with the first sub-tokens, ignoring *pad* tokens in the loss function. Figure 1 illustrates this alignment. We experiment with *BERT-base* and *BERT-large* models, exploring full finetuning, partial finetuning of the lower layers, and finetuning only the last layer.

In both experiments, training proceeds for a maximum of 200 epochs, with early stopping triggered after 5 consecutive validation stagnations. We use *AdamW* as the optimizer and clip gradients to ensure stability. A larger learning rate ($1e-4$) is applied to the classifiers, while a smaller rate ($5e-5$) is used to fine-tune *BERT*. Configuration details are provided in Table 2.

Intent Classification and *Slot Filling* tasks employ the *Cross Entropy* loss function, excluding *<pad>* tokens. The losses are summed and minimized jointly. Intent classification is evaluated using accuracy, while slot filling is measured using the *CoNLL* script to report the F1 score.

We use the *ATIS* dataset for training and evaluation. For experiments involving the *LSTM* model, a vocabulary is created to map words to IDs. For *BERT*-based models, we employ the *BERT-tokenizer*. Specifically, we consider *BERT-base* and *BERT-large* as backbones. Parameters for the data loaders are summarized in Table 3.

3. Results

In all experiments, we report the average score over 5 runs along with the relative standard deviation. Results from both exper-

iments are summarized in Table 4 and illustrated in Figures 2 and 3.

Our baseline achieves 93.5% on Intent Classification and 91.1% on Slot Filling. Adding bidirectionality results in a $> 1.5\%$ improvement in both tasks, while introducing dropout has minimal impact on performance, though it reduces the standard deviation. Figures 4 and 5 show the progression of intent accuracy and slot F1 scores on the validation set.

Using *BERT* as the backbone significantly enhances performance, with the best settings achieving 97.5% accuracy on Intent Classification and 96.1% on Slot Filling. Notably, fine-tuning only the last layer surpasses the performance of the *LSTM*-based model in both tasks. Results are comparable between *BERT-base* and *BERT-large*, with the best performance obtained by fine-tuning the entire model. Figures 4 and 5 further depict the trends in intent accuracy and slot F1 scores on the validation set.

4. References

- [1] Q. Chen, Z. Zhuo, and W. Wang, “Bert for joint intent classification and slot filling,” 2019. [Online]. Available: <https://arxiv.org/abs/1902.10909>

A. Tables and Figures

LSTM hyperparameters	
Embedding size	300
Hidden size	200
Slot Classifier	130
Intent Classifier	26
n layers	1
pad index	0
Output dropout	0.1
Embedding dropout	0.1

Table 1: LSTM hyperparameters

Training hyperparameters	
n epochs	200
n epochs for fine-tuning	50
patience	5
clip	5
training loss	Cross Entropy
Intent evaluation metric	Accuracy
Slot filling metric	F1 score
optimizer	AdamW
learning rate	1 e-4
learning rate bert parameters	5 e-5
weight decay	0.01

Table 2: Training hyperparameters

Training loader	
Length training dataset	4480
Batchsize	128
Number of batches	35
Validation loader	
Length validation dataset	498
Batchsize	64
Number of batches	8
Test loader	
Length test dataset	893
Batchsize	64
Number of batches	14

Table 3: Loader hyperparameters

Model	Intent		Slot filling	
	Accuracy	std	F1-score	std
LSTM	93.5	0.6	92.1	0.4
LSTM bidirectional	94.8	0.5	93.7	0.2
LSTM bidir + dropout	94.7	0.1	93.9	0.1
Bert_base fully tuned	97.5	0.3	95.7	0.1
Bert_base partially tuned	97.4	0.3	95.4	0.1
Bert_base 1 layer	97.2	0.1	94.0	0.9
Bert_large fully tuned	97.4	0.2	96.1	0.2
Bert_large partially tuned	97.4	0.2	95.7	0.1
Bert_large 1 layer	96.7	0.2	91.9	0.3

Table 4: NLU performance on ATIS dataset. The metrics are intent classification accuracy and slot filling F1 (%). Results are averaged over 5 runs with relative standard deviation (std).

Original Tokens	First sub-token	Slot Labels	Slot IDs
[CLS]	---	<pad>	0
show	show	0	43
me	me	0	43
the	the	0	43
cheap	cheap	B-cost_relative	35
##est	---	<pad>	0
flight	flight	0	43
from	from	0	43
pittsburgh	pittsburgh	B-fromloc.city_name	46
to	to	0	43
atlanta	atlanta	B-toloc.city_name	12
on	on	0	43
wednesday	wednesday	B-depart_date.day_name	74
which	which	0	43
leaves	leaves	0	43
before	before	B-depart_time.time_relative	117
noon	noon	B-depart_time.time	37
and	and	0	43
serves	serves	0	43
breakfast	breakfast	B-meal_description	76
[SEP]	---	<pad>	0

Figure 1: Tokenized utterance aligned with slot labels.




Figure 2: Comprehensive comparison




Figure 3: Comprehensive comparison




Figure 4: *Intent Accuracy on validation set in experiment 1*




Figure 5: *Slot F1 on validation set in experiment 1*




Figure 6: *Intent Accuracy on validation set in experiment 2*




Figure 7: *Slot F1 on validation set in experiment 2*